

# Django API Optimization Guide

## Performance Impact Summary

### Immediate Impact (50-90% improvement)

- **Prefetching:** Reduces 100+ queries to 2-3 queries
- **Database Indexes:** 10-100x faster filtering and searching
- **Caching Fields:** Eliminates expensive COUNT() operations

### Scalability Impact

- **Soft Delete:** Maintains performance as data grows
- **Denormalization:** Dashboard queries stay fast regardless of historical data
- **Custom QuerySets:** Consistent optimization patterns across your API

### User Experience Impact

- **Search Optimization:** Sub-second search results even with thousands of records
- **Pagination:** Large datasets load smoothly
- **Response Caching:** Frequently accessed data serves instantly

## Implementation Priority

1. **Start with:** Database indexes and prefetching (biggest immediate impact)
  2. **Then add:** Custom QuerySets and caching fields
  3. **Advanced:** Full-text search and denormalization for complex analytics
- 

## 4. Soft Delete Pattern

### What it is:

Instead of permanently deleting records, mark them as deleted with a timestamp.

python

```
class BaseModel(models.Model):
    created_at = models.DateTimeField(auto_now_add=True)
    updated_at = models.DateTimeField(auto_now=True)
    deleted_at = models.DateTimeField(null=True, blank=True)

    objects = models.Manager() # All records
    active_objects = ActiveManager() # Only non-deleted records

    class Meta:
        abstract = True

    def soft_delete(self):
        self.deleted_at = timezone.now()
        self.save()

    def is_deleted(self):
        return self.deleted_at is not None

class ActiveManager(models.Manager):
    def get_queryset(self):
        return super().get_queryset().filter(deleted_at__isnull=True)
```

## API Benefits:

- **Data Recovery:** Can restore "deleted" records
- **Audit Trail:** Keep history of what was deleted and when
- **Referential Integrity:** Related records don't break when something is "deleted"
- **Better Analytics:** Historical data remains for reporting

## Usage in APIs:

python

```
# Instead of hard delete
employee.delete() # Permanent deletion

# Use soft delete
employee.soft_delete() # Recoverable deletion

# Query active records
active_employees = Employee.active_objects.all()
```

## 5. Caching Fields for Performance

### What it is:

Store computed values in the database to avoid expensive calculations on every API call.

```
python

class Branch(models.Model):
    # ... existing fields ...
    employee_count = models.PositiveIntegerField(default=0)
    active_employee_count = models.PositiveIntegerField(default=0)
    last_meal_log_date = models.DateTimeField(null=True, blank=True)

    def update_counts(self):
        self.employee_count = self.branch_employees.count()
        self.active_employee_count = self.branch_employees.filter(is_active=True).count()
        self.save(update_fields=['employee_count', 'active_employee_count'])

class Company(models.Model):
    total_branches = models.PositiveIntegerField(default=0)
    total_employees = models.PositiveIntegerField(default=0)
    active_branches = models.PositiveIntegerField(default=0)
```

## API Benefits:

- **Faster Responses:** No need to count records on every request
- **Reduced Database Load:** Fewer complex queries
- **Consistent Performance:** Response time doesn't increase with data growth

## Update Strategies:

```
python

# Using Django signals
from django.db.models.signals import post_save, post_delete

@receiver(post_save, sender=Employee)
def update_branch_counts(sender, instance, **kwargs):
    if instance.branch:
        instance.branch.update_counts()

# Or periodic tasks with Celery
@periodic_task(run_every=crontab(minute=0)) # Every hour
def update_all_counts():
    for branch in Branch.objects.all():
        branch.update_counts()
```

---

## 6. Custom QuerySets for Query Optimization

### What it is:

Pre-defined query methods that include optimal database operations.

python

```
class EmployeeQuerySet(models.QuerySet):
    def with_related(self):
        """Load related data in single query"""
        return self.select_related(
            'branch', 'department', 'employee_type', 'canteen'
        ).prefetch_related('employee_meal_logs')

    def active(self):
        return self.filter(is_active=True, deleted_at__isnull=True)

    def by_branch(self, branch_id):
        return self.filter(branch_id=branch_id)

    def with_meal_summary(self):
        """Add meal statistics"""
        return self.annotate(
            total_meals=Count('employee_meal_logs'),
            meals_this_month=Count(
                'employee_meal_logs',
                filter=Q(employee_meal_logs__created_at__month=timezone.now().month)
            )
        )

    def search(self, query):
        """Full-text search across multiple fields"""
        return self.filter(
            Q(name__icontains=query) |
            Q(employee_id__icontains=query) |
            Q(branch__name__icontains=query)
        )

class Employee(models.Model):
    # ... fields ...
    objects = EmployeeQuerySet.as_manager()
```

## API Benefits:

- **Consistent Queries:** Same optimization patterns across your app
- **Reduced N+1 Problems:** Built-in related data loading
- **Maintainable Code:** Complex queries defined once, used everywhere

## Usage in API Views:

```
python
```

```
# Bad: Multiple database hits
```

```
def get_employees(request):  
    employees = Employee.objects.all()  
    # Each employee access triggers additional queries for branch, department, etc.
```

```
# Good: Single optimized query
```

```
def get_employees(request):  
    employees = Employee.objects.with_related().active()  
    # All data loaded in one query
```

## 9. Prefetch Strategies

### What it is:

Loading related data efficiently to avoid N+1 query problems.

```
python
```

```
# Problem: N+1 queries
```

```
employees = Employee.objects.all()  
for employee in employees: # 1 query  
    print(employee.branch.name) # N additional queries
```

```
# Solution: Prefetching
```

```
employees = Employee.objects.select_related('branch', 'department')  
for employee in employees: # 1 query total  
    print(employee.branch.name) # No additional queries
```

### Types of Prefetching:

#### select\_related (for ForeignKey and OneToOne):

```
python
```

```
# Loads related data in same query using JOINS
```

```
employees = Employee.objects.select_related(  
    'branch',  
    'department',  
    'canteen',  
    'branch__company' # Can chain relationships  
)
```

#### prefetch\_related (for ManyToMany and reverse ForeignKey):

```
python
```

```
# Loads related data in separate optimized queries
```

```
branches = Branch.objects.prefetch_related(  
    'branch_employees', # All employees for each branch  
    'branch_departments', # All departments for each branch  
    'branch_employees__employee_meal_logs' # Meal logs for all employees  
)
```

## Advanced Prefetching:

```
python
```

```
from django.db.models import Prefetch
```

```
# Custom prefetch with filtering
```

```
recent_logs = MealLog.objects.filter(  
    created_at__gte=timezone.now() - timedelta(days=30)  
)
```

```
employees = Employee.objects.prefetch_related(  
    Prefetch('employee_meal_logs', queryset=recent_logs, to_attr='recent_meals')  
)
```

```
# Now employee.recent_meals contains only last 30 days
```

## API Performance Impact:

```
python
```

```
# Before optimization: 1 + N queries
```

```
# GET /api/employees/ with 100 employees = 101 database queries
```

```
# After optimization: 2-3 queries total
```

```
employees = Employee.objects.select_related('branch', 'department').prefetch_related('employee_meal_logs')
```

---

## 11. Search Fields and Full-Text Search

### What it is:

Optimized search functionality for API endpoints.

### Basic Search Implementation:

python

```
class EmployeeQuerySet(models.QuerySet):
    def search(self, query):
        if not query:
            return self

        return self.filter(
            Q(name__icontains=query) |
            Q(employee_id__icontains=query) |
            Q(branch__name__icontains=query) |
            Q(department__name__icontains=query)
        )
```

## PostgreSQL Full-Text Search:

python

```
from django.contrib.postgres.search import SearchVector, SearchQuery, SearchRank
from django.contrib.postgres.indexes import GinIndex

class Employee(models.Model):
    # ... existing fields ...
    search_vector = SearchVectorField(null=True)

    class Meta:
        indexes = [
            GinIndex(fields=['search_vector']),
        ]

    # Update search vector
    Employee.objects.update(
        search_vector=SearchVector('name', 'employee_id', 'branch__name')
    )

    # Search usage
    def search_employees(query):
        search_query = SearchQuery(query)
        return Employee.objects.annotate(
            rank=SearchRank('search_vector', search_query)
        ).filter(search_vector=search_query).order_by('-rank')
```

## API Benefits:

- **Fast Search:** Indexed search is much faster than LIKE queries
- **Relevance Ranking:** Results ordered by relevance
- **Multi-field Search:** Search across multiple fields simultaneously

## 12. Denormalization for Frequent Queries

### What it is:

Storing computed/aggregated data in separate tables for faster access.

python

```
class EmployeeDashboard(models.Model):
    """Denormalized data for dashboard APIs"""
    employee = models.OneToOneField(Employee, on_delete=models.CASCADE)

    # Cached company/branch info
    company_name = models.CharField(max_length=100)
    branch_name = models.CharField(max_length=100)
    department_name = models.CharField(max_length=100)

    # Meal statistics
    total_meals_current_month = models.PositiveIntegerField(default=0)
    total_meals_last_month = models.PositiveIntegerField(default=0)
    favorite_meal_type = models.CharField(max_length=20, null=True)
    last_meal_date = models.DateTimeField(null=True)

    # Cost information
    total_meal_cost_current_month = models.DecimalField(max_digits=10, decimal_places=2, default=0)

    updated_at = models.DateTimeField(auto_now=True)

class CompanyStats(models.Model):
    """Company-level statistics"""
    company = models.OneToOneField(Company, on_delete=models.CASCADE)
    total_employees = models.PositiveIntegerField(default=0)
    active_employees = models.PositiveIntegerField(default=0)
    total_branches = models.PositiveIntegerField(default=0)
    meals_served_today = models.PositiveIntegerField(default=0)
    meals_served_this_month = models.PositiveIntegerField(default=0)
    top_meal_type = models.CharField(max_length=20, null=True)
    updated_at = models.DateTimeField(auto_now=True)
```

### Update Strategy:



```
python
```

```
# Real-time updates via signals
```

```
@receiver(post_save, sender=MealLog)
```

```
def update_employee_dashboard(sender, instance, created, **kwargs):
```

```
    if created:
```

```
        dashboard, _ = EmployeeDashboard.objects.get_or_create(
```

```
            employee=instance.employee
```

```
        )
```

```
        dashboard.refresh_stats()
```

```
# Periodic batch updates
```

```
@periodic_task(run_every=crontab(minute=0)) # Every hour
```

```
def refresh_dashboard_stats():
```

```
    for dashboard in EmployeeDashboard.objects.all():
```

```
        dashboard.refresh_stats()
```

## API Benefits:

- **Lightning Fast Dashboards:** Pre-computed data loads instantly
- **Complex Analytics:** Expensive calculations done offline
- **Consistent Performance:** Response time independent of historical data size

---

## Additional Performance Recommendations

### Database Connection Optimization:

```
python
```

```
# settings.py
```

```
DATABASES = {
```

```
    'default': {
```

```
        'ENGINE': 'django.db.backends.postgresql',
```

```
        'CONN_MAX_AGE': 600, # Persistent connections
```

```
        'OPTIONS': {
```

```
            'MAX_CONNS': 20, # Connection pooling
```

```
        }
```

```
    }
```

```
}
```

### Query Debugging:

```
python
```

```
# settings.py (development)
```

```
LOGGING = {  
    'loggers': {  
        'django.db.backends': {  
            'level': 'DEBUG', # Log all SQL queries  
        }  
    }  
}
```

```
# Use django-debug-toolbar
```

```
INSTALLED_APPS = [  
    'debug_toolbar',  
]
```

## API Response Caching:

```
python
```

```
from django.views.decorators.cache import cache_page  
from django.core.cache import cache
```

```
@cache_page(60 * 15) # Cache for 15 minutes
```

```
def employee_list(request):  
    return JsonResponse(employee_data)
```

```
# Or manual caching
```

```
def get_company_stats(company_id):  
    cache_key = f'company_stats_{company_id}'  
    stats = cache.get(cache_key)  
  
    if stats is None:  
        stats = calculate_company_stats(company_id)  
        cache.set(cache_key, stats, 60 * 60) # Cache 1 hour  
  
    return stats
```

## Pagination for Large Datasets:

python

```
from django.core.paginator import Paginator
from rest_framework.pagination import PageNumberPagination

class StandardResultsSetPagination(PageNumberPagination):
    page_size = 50
    page_size_query_param = 'page_size'
    max_page_size = 1000

# Cursor pagination for very large datasets
from rest_framework.pagination import CursorPagination

class EmployeeCursorPagination(CursorPagination):
    page_size = 50
    ordering = 'id' # Must have consistent ordering
```

These optimizations can dramatically improve API performance, reducing response times from seconds to milliseconds and supporting much larger datasets efficiently.